CSSE 304   Assignment #13 (interpreter assignment #1) Updated for Spring, 2016

This is a **pair assignment** (possibly with a few solo or trio groups). The intention is not that you will use a divide-and-conquer approach, but that you will instead get together to do the work (i.e., pair programming) so that everyone is involved in the solutions of all parts, and everyone understands all of it. I will expect everyone to understand all of the code that your pair submits, and that you will do your best to make sure that your partner understands it. For each interpreter assignment, there will be a brief survey on Moodle in which you will be asked verify that this is the case.

When you submit your code
- Include the usernames of other group member(s) on the PLC server's submission page (only necessary for your submission that is to be graded).
  If your interpreter code is distributed over several files (I am giving you both single-file and multi-file starting code):
  a. Your main file should be `main.ss`. It should load the other files. The argument to each `load` statement should be just a filename, not a full pathname.
  b. Create a `.zip` file containing all of your source files, including `main.ss`. You do not have to include `chez-init.ss` in your `.zip` file., but it is okay if you do.
  c. Submit this `.zip` file to the A13 assignment on the server.
- After your team's final submission, each member must fill out the participation survey on Moodle (unless you are a "solo team").

Late day usage: It is possible that (on this or a future interpreter assignment) one partner (A) will run out of late days while the other partner (B) still has late days left. If this happens, partner A's score will be based on the best submission before the due date, while Partner B's score will be based on the best submission before the late day deadline.   If you are a potential Partner B in this situation, the best thing you can do for both of you is to push your partner to work with you early so you can submit an assignment early so that partner A will have a late day to use when you both need it.

Collaborative test case development: If you think you have good test(s) that cover things (or combinations of things) that my test cases don't cover, please post them in the *interpreter_project* folder on Piazza. I may add some of your test cases to the ones that the grading program uses. If I use your test case(s), I will give you a few bonus homework points. Even if I don't use your new tests, it will be good for everyone to have a richer test bed.

(135 points) programming problem
This assignment is not conceptually difficult, but there are many details to understand and to implement. The simple parser and interpreter presented in class can be obtained from
- http://www.rose-hulman.edu/class/csse/csse304/201610/Homework/Assignment_13/A13-multiple-files.zip (a separate file for each major part of the code, perhaps easiest for the development process), or
- http://www.rose-hulman.edu/class/csse/csse304/201610/Homework/Assignment_13/all.ss (a single file containing all of the starting code, easiest to submit, because you do not have to make a ZIP file each time)

Your first job is to understand the given code, and then you will add features to the given parser and interpreter. You should substitute part or all of your own A11 parser for the simple parser that I provide. Some of the code that you need to add to your interpreter is in the EoPL book. Don't just blindly add it to your code; instead be sure that you understand everything as you are adding it.

Interpreted language: The textbook describes an alternate, Pascal-like syntax for the interpreted language. The authors' rationale is that you are less likely to get confused between the interpreted language and the implementation language (Scheme) in which you write the interpreter. However, I think you can handle interpreting Scheme-like syntax without much confusion, and that the benefits of writing a Scheme-like

interpreter outweigh the disadvantages. For one thing, you can be sure that your interpreter is getting the correct answer for most expressions, because you can simply input the same expressions directly into Petite *Chez* Scheme and see what Scheme returns. Thus, whenever an exercise in the book describes a different syntax, you should use standard Scheme syntax instead.

Summary: The major features that you are to add to the interpreted language in this assignment are:
- additional primitive procedures
- literals
  - Boolean constants `#t`, and `#f`
  - quoted values, such as `'( )`, `'(a b c)`, `'(a b . (c))`, `'#(2 5 4)`
  - string literals, such as `"abc"`
- `if` two-armed only: i.e., `(if test-exp then-exp else-exp)`
- `let` (not `let*` or named `let` yet; those will appear in later assignments; don't remove them from your parser, though)
- `lambda` (just the normal `lambda` with a proper list of arguments)
- `rep` and `eval-one-exp` (two alternative interfaces to your interpreter, described below)

- I suggest that you thoroughly test each feature before adding the next one. Augmenting `unparsed-exp` whenever you augment `parse-exp` is a good idea, to help you with debugging.

  - Several of my test cases will involve `let` and/or `lambda`.

A more detailed description of what you are to do:

For each feature in this and subsequent interpreter assignments, unless specified otherwise, the syntax and semantics should be the same as Scheme's.

Add the Boolean constants `#t`, and `#f` to the interpreted language, along with quoted data, such as lists and vectors. Also add string literals such as `"abcd"`. You do not need to add any string-manipulation procedures yet, but it will be nice to at least be able to use strings for output messages. Add vector literals (as with strings, Scheme will do the "real" parsing; your `parse-exp` simply needs to call `vector?` to see if an expression is a vector).

Add `if` expressions to the interpreted language. Most of the code is in the textbook, but you will have to adapt it to recognize and use Boolean and other literal values— in addition to the numeric values. In the book's interpreted language, the authors represent *true* and *false* by numbers. In Scheme (and thus in your interpreter), any non-*false* value should be treated as *true*. The number 0 is a *true* value in Scheme (and in your interpreted language), but 0 is the false value in the textbook's language).
Note: Recall that `if` has two forms, with and without an "else" expression. Your interpreter must eventually support both, but only the "with" version is required for this assignment.

Add several primitive procedures including `+`, `-`, `*`, `/`, `add1`, `sub1`, `zero?`, `not`, `=` and `<` (and the other numeric comparison operators), and also `cons`, `car`, `cdr`, `list`, `null?`, `assq`, `eq?`, `equal?`, `atom?`, `length`, `list->vector`, `list?`, `pair?`, `procedure?`, `vector->list`, `vector`, `make-vector`, `vector-ref`, `vector?`, `number?`, `symbol?`, `set-car!`, `set-cdr!`, `vector-set!`, `display`, `newline` to your interpreter. Add the `c**r` and `c***r` procedures (where each "`*`" stands for an "`a`" or "`d`").
Note: You may use built-in Scheme procedures in your implementation of most of the primitive procedures, as I do in the starting code that I provided for you. At least one of the above procedures has a much simpler implementation, and at least one cannot be implemented by using the corresponding Scheme procedure. You may want to enhance the `prim-proc` mechanism from the provided code so

that a `prim-proc` knows how many arguments it can take, and reports an error if the interpreted code attempts to apply it with an incorrect number of arguments. Otherwise, this incorrect code will default to a Scheme error message such as "Exception in cadr", which will make no sense to the author of the interpreted code.

Add `let` (not named-`let`) expressions to the parser, with the standard Scheme syntax.
`(let ([var exp]…) body1 body2 …)`. The `...` is basically the same thing as the Kleene * (like the notation we used with `define-syntax`); it means "zero or more occurrences of the previous item". Every `let` expression must have at least one body. (This work should have been done in a previous assignment). It is okay for you to interpret named `let` as well, but this is not required until a later assignment.

Add code similar to that in EoPL 3.2, allowing `let` expressions to be directly interpreted. Note that after the extended environment is created, the bodies are executed in order, and the value of the last body is returned. I suggest that you test this code and understand it thoroughly before you go on to implement `lambda`.

Add `lambda` expressions of the form `(lambda (var ...) body1 body2 ...)`, See the description of `rep` (short for "read-eval-print") below for a description of what to print if a closure is part of the final value of an expression entered interactively in the interpreter. Section 3.3 of EoPL may be helpful.

In later assignments, you'll add many syntactic and semantic forms, including `define`, `let*`, `letrec`, and `set!`

INTERFACES TO YOUR INTERPRETER

A. A normal interactive interface (read-eval-print loop)

This interface will be used primarily for your interactive testing of your interpreter and to facilitate interaction if you bring the interpreter to me for help. The user should load your code, type
`(rep)`
and begin entering Scheme expressions. Your read-eval-print loop should display a prompt that is different from the normal Scheme prompt. It should read and evaluate an expression, print the value, and then prompt for the next expression. If the value happens to be (or contain) a closure, your interpreter should not print the internal representation of the closure, but instead it should print `<interpreter-procedure>`. For debugging purposes, you might want to write a separate procedure called `rep-debug` that behaves like `rep`, except that it displays the internal representation of everything that it prints, including procedures. Using `trace` on your `eval-exp` procedure can be a tremendous help, but be prepared to wade through a lot of output!

If the user types `(exit)` as an input expression to your interpreter, (some later version of) your interpreter should exit and simply return to the Scheme top-level. What happens when the user enters illegal Scheme code or encounters a run-time error? Ideally, the interpreter should print an error message and return to the "read" portion of the "read-eval-print" loop. Think about how you might do that. But I do not expect you to do it for this assignment. A sample transcript follows:

```
> (rep)
--> (+ 3 5)
8
--> (list 6 (lambda (x) x) 7)
(6 <interpreter-procedure> 7)
--> (cons 'a '(b c d))
(a b c d)
--> (let ([a 4])
      (+ a 6))
10
--> ((lambda (x)
      ((lambda (y)
        (+ x y))
       5))
     6)
11
--> (let ([t '(3 4)])
      (if (< (car t) 1)
          7
          (let ()
            (set-car! t (+ 1 (cadr t)))
            (cons (cadr t) (car t)))))
(4 . 5)
--> (((lambda (f)
        ((lambda (x) (f (lambda (y) ((x x) y))))
         (lambda (x) (f (lambda (y) ((x x) y))))))
      (lambda (g)
        (lambda (n)
          (if (zero? n)
              1
              (* n (g (- n 1)))))))
   6)
720
-->(exit)
```

B. single-procedure interface: `eval-one-exp`

In addition to the (rep) interactive interface, you must provide the procedure `eval-one-exp`. The code I have given you below is intended to clarify its function, not to make you rewrite your interpreter. You will of course need to adapt it to your particular program. It may be possible to arrange things so that your rep procedure can call `eval-one-exp`.

```
(define eval-one-exp ; may need to be enhanced.
  (lambda (exp) (eval-exp (parse-exp exp))))))
```

```
>(eval-one-exp '(- (* 2 3) (* 6 3)))
-12
>
```

Note that `eval-one-exp` **does not have to be available to the user when running your interpreter interactively.** It only has to be available from the normal Scheme prompt.